

# Announcements

---

- **Project 0: Python Tutorial**
  - Due tomorrow!
  - There is a lab **Wednesday from 3pm-5pm in Soda 275**
  - The lab time is optional, but P0 itself is not
  - On submit, you should get email from the autograder
- **Project 1: Search**
  - On the web today
  - Start early and ask questions. It's longer than most!
- **Self-Diagnostic on web**
- **Sections: can go to any, but have priority in your own**

## CS 188: Artificial Intelligence Fall 2011

### Lecture 2: Queue-Based Search 8/30/2011

Dan Klein – UC Berkeley

Multiple slides from Stuart Russell, Andrew Moore

# Today



---

- Agents that Plan Ahead
- Search Problems
- Uninformed Search Methods (part review for some)
  - Depth-First Search
  - Breadth-First Search
  - Uniform-Cost Search
- Heuristic Search Methods (new for all)
  - Greedy Search

## Reflex Agents

---

- Reflex agents:
  - Choose action based on current percept (and maybe memory)
  - May have memory or a model of the world's current state
  - Do not consider the future consequences of their actions
  - Consider how the world IS
- Can a reflex agent be rational?





[demo: reflex optimal / loop ]

# Goal Based Agents

- Goal-based agents:

- Plan ahead
- Ask “what if”
- Decisions based on (hypothesized) consequences of actions
- Must have a model of how the world evolves in response to actions
- Consider how the world **WOULD BE**



[demo: plan fast / slow ]


# Search Problems

- A **search problem** consists of:

- A state space



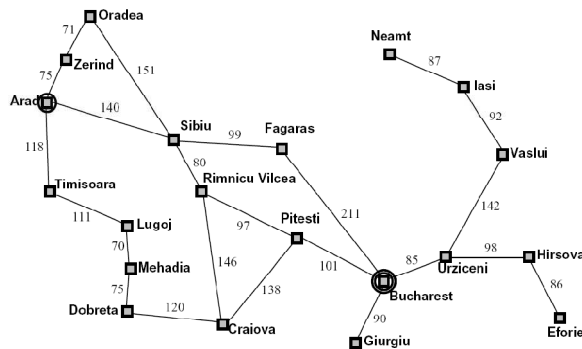
- A successor function (with actions, costs)



- A start state and a goal test

- A **solution** is a sequence of actions (a plan) which transforms the start state to a goal state


## Example: Romania



- **State space:**
  - Cities
- **Successor function:**
  - Roads: Go to adj city with cost = dist
- **Start state:**
  - Arad
- **Goal test:**
  - Is state == Bucharest?
- **Solution?**

## State Space Graphs


- **State space graph: A mathematical representation of a search problem**
  - For every search problem, there's a corresponding state space graph
  - The successor function is represented by arcs
- **We can rarely build this graph in memory (so we don't)**



*Ridiculously tiny search graph for a tiny search problem*

# What's in a State Space?

The **world state** specifies every last detail of the environment



A **search state** keeps only the details needed (abstraction)

## Problem: Pathing

- States: (x,y) location
- Actions: NSEW
- Successor: update location only
- Goal test: is (x,y)=END

## Problem: Eat-All-Dots

- States: {(x,y), dot booleans}
- Actions: NSEW
- Successor: update location and possibly a dot boolean
- Goal test: dots all false


# State Space Sizes?

## World state:


- Agent positions: 120
- Food count: 30
- Ghost positions: 12
- Agent facing: NSEW

## How many

- World states?  
 $120 \times (2^{30}) \times (12^2) \times 4$
- States for pathing?  
120
- States for eat-all-dots?  
 $120 \times (2^{30})$




# Search Trees



- **A search tree:**

- This is a “what if” tree of plans and outcomes
- Start state at the root node
- Children correspond to successors
- Nodes contain states, correspond to PLANS to those states
- For most problems, we can never actually build the whole tree

## Another Search Tree



- **Search:**

- Expand out possible plans
- Maintain a **fringe** of unexpanded plans
- Try to expand as few tree nodes as possible

# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```


- Important ideas:

- Fringe
- Expansion
- Exploration strategy


*Detailed pseudocode  
is in the book!*

- Main question: which fringe nodes to explore?

## Example: Tree Search




# State Graphs vs. Search Trees



Each NODE in the search tree is an entire PATH in the problem graph.



We construct both on demand – and we construct as little as possible.



# Review: Depth First Search

Strategy: expand deepest node first

Implementation: Fringe is a LIFO stack







# Review: Breadth First Search

Strategy: expand shallowest node first

Implementation:  
Fringe is a FIFO queue



# Search Algorithm Properties

**Complete?** Guaranteed to find a solution if one exists?

**Optimal?** Guaranteed to find the least cost path?

**Time complexity?**


**Space complexity?**

Variables:

|       |  |
|-------|--|
| $n$   | Number of states in the problem (huge)                                 |
| $b$   | The average branching factor $B$<br>(the average number of successors) |
| $C^*$ | Cost of least cost solution  |
| $s$   | Depth of the shallowest solution                                       |
| $m$   | Max depth of the search tree   |

# DFS


| Algorithm |                    | Complete | Optimal | Time     | Space    |
|-----------|--------------------|----------|---------|----------|----------|
| DFS       | Depth First Search | N        | N       | Infinite | Infinite |



- Infinite paths make DFS incomplete...
- How can we fix this?

# DFS

- With cycle checking, DFS is complete.\*




| Algorithm |                  | Complete | Optimal | Time         | Space   |
|-----------|------------------|----------|---------|--------------|---------|
| DFS       | w/ Path Checking | Y        | N       | $O(b^{m+1})$ | $O(bm)$ |

- When is DFS optimal?

\* Or graph search – next lecture.

# BFS

| Algorithm |                  | Complete | Optimal | Time         | Space    |
|-----------|------------------|----------|---------|--------------|----------|
| DFS       | w/ Path Checking | Y        | N       | $O(b^{m+1})$ | $O(bm)$  |
| BFS       |                  | Y        | N*      | $O(b^{s+1})$ | $O(b^s)$ |



- When is BFS optimal?


# Comparisons

- When will BFS outperform DFS?
- When will DFS outperform BFS?

# Iterative Deepening


Iterative deepening: BFS using DFS as a subroutine:

1. Do a DFS which only searches for paths of length 1 or less.
2. If "1" failed, do a DFS which only searches paths of length 2 or less.
3. If "2" failed, do a DFS which only searches paths of length 3 or less.  
....and so on.



| Algorithm |                  | Complete | Optimal | Time         | Space    |
|-----------|------------------|----------|---------|--------------|----------|
| DFS       | w/ Path Checking | Y        | N       | $O(b^{m+1})$ | $O(bm)$  |
| BFS       |                  | Y        | N*      | $O(b^{s+1})$ | $O(b^s)$ |
| ID        |                  | Y        | N*      | $O(b^{s+1})$ | $O(bs)$  |

# Costs on Actions





Notice that BFS finds the shortest path in terms of number of transitions. It does not find the least-cost path.  
We will quickly cover an algorithm which does find the least-cost path.

# Uniform Cost Search

Expand cheapest node first:

Fringe is a priority queue  
(priority: cumulative cost)



## Priority Queue Refresher


- A priority queue is a data structure in which you can insert and retrieve (key, value) pairs with the following operations:

|                     |   |
|---------------------|---|
| pq.push(key, value) | inserts (key, value) into the queue.                                  |
| pq.pop()            | returns the key with the lowest value, and removes it from the queue. |

- You can decrease a key's priority by pushing it again
- Unlike a regular queue, insertions aren't constant time, usually  $O(\log n)$
- We'll need priority queues for cost-sensitive search methods

# Uniform Cost Search



| Algorithm |                  | Complete | Optimal | Time                  | Space                 |
|-----------|------------------|----------|---------|-----------------------|-----------------------|
| DFS       | w/ Path Checking | Y        | N       | $O(b^{m+1})$          | $O(bm)$               |
| BFS       |                  | Y        | N       | $O(b^{s+1})$          | $O(b^s)$              |
| UCS       |                  | Y*       | Y       | $O(b^{C^*/\epsilon})$ | $O(b^{C^*/\epsilon})$ |



\* UCS can fail if actions can get arbitrarily cheap

# Uniform Cost Issues


- Remember: explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
  - Explores options in every "direction"
  - No information about goal location



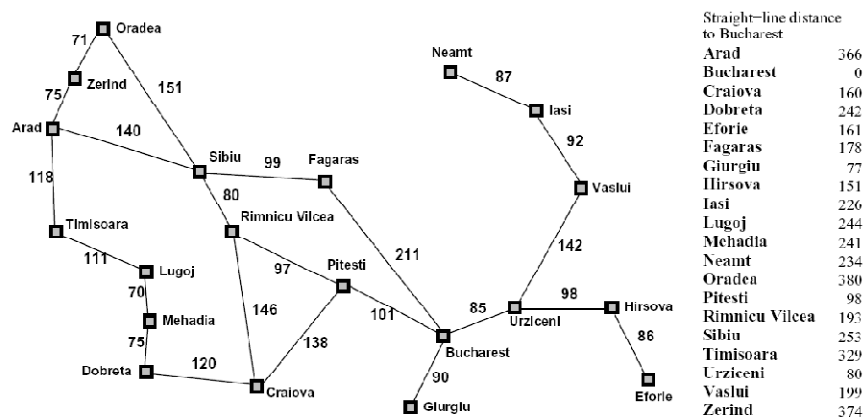
[demo: search demo empty]

# Search Heuristics

- Any *estimate* of how close a state is to a goal
- Designed for a particular search problem
- Examples: Manhattan distance, Euclidean distance

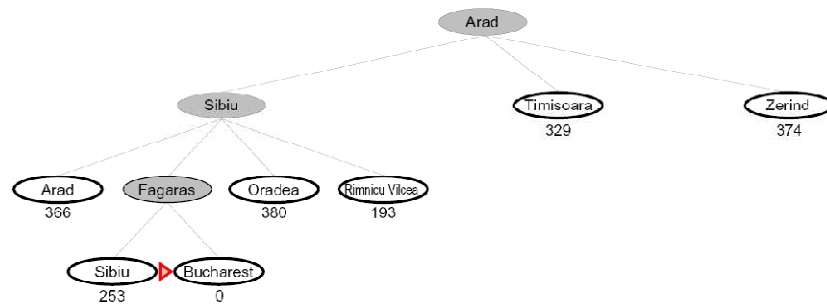


# Heuristics



## Best First / Greedy Search

- Expand the node that seems closest...





- What can go wrong?

[demo: greedy]

## Best First / Greedy Search

- A common case:
  - Best-first takes you straight to the (wrong) goal
- Worst-case: like a badly-guided DFS in the worst case
  - Can explore everything
  - Can get stuck in loops if no cycle checking
- Like DFS in completeness (finite states w/ cycle checking)





# Search Gone Wrong?

